

STA355 Homework 3

Luis Rojas

21/03/2021

(c) The file `bees.txt` contains dance directions of 279 honey bees viewing a zenith patch of artificially polarized light. (The data are given in degrees; you should convert them to radians. The original data were rounded to the nearest 10° ; the data in the file have been “jittered” to make them less discrete.) Using these data, compute the posterior density of κ in part (b) for $\lambda = 1$ and $\lambda = 0.1$. How do these two posterior densities differ? Do these posterior densities “rule out” the possibility that $\kappa = 0$ (i.e. a uniform distribution)?

Note: To compute the posterior density, you will need to compute the normalizing constant – on Quercus, I will give some suggestions on how to do this. A simple estimate of κ is

$$\hat{\kappa} = \frac{\frac{r}{n} \left(2 - \frac{r^2}{n^2} \right)}{1 - \frac{r^2}{n^2}}$$

where r is as defined in part (b). The posterior density should be largest for values of κ close to $\hat{\kappa}$; you can use this as a guide to determine for what range of values of κ you should compute the posterior density.

```
bees <- scan("bees.txt")
```

(The data are given in degrees; you should convert them to radians. The original data were rounded to the nearest 10degree ; the data in the file have been “jittered” to make them less discrete.)

Using these data, compute the posterior density of κ in part (b) for $\lambda = 1$ and $\lambda = 0.1$. How do these two posterior densities differ? Do these posterior densities “rule out” the possibility that $\kappa = 0$ (i.e. a uniform distribution)?

From b we have:

$$\pi(k|d_1, \dots, d_n) = c(d_1, \dots, d_n) \frac{\exp(-\alpha\kappa) I_0(r\kappa)}{[I_0(\kappa)]^n}$$

The von Misses Distribution whose density on $[0, 2\pi]$ is:

$$f(\theta; \kappa, \mu) = \frac{1}{2\pi I_0(\kappa)} \exp(\kappa \cos(\theta - \mu))$$

We know the likelihood (considering only κ) has the form:

$$L(D_i ; \kappa, \mu) = \left(\frac{1}{2\pi} \right)^n \left(\frac{1}{\prod_{i=1}^n I_0(\kappa)} \right) \int_0^{2\pi} \exp \left[\sum_{i=1}^n \kappa \cos(D_i - \mu) \right] d\mu$$

Writing in terms of the Bessel function of the first kind:

$$L(D_i ; \kappa, \mu) = \frac{1}{(2\pi)^{n-1}} \frac{1}{\prod_{i=1}^n I_0(\kappa)} \frac{1}{2\pi} \int_0^{2\pi} \exp [\kappa r \cos(\theta - \mu)] d\mu$$

$$L(D_i ; \kappa, \mu) = \frac{1}{(2\pi)^{n-1}} \frac{1}{\prod_{i=1}^n I_0(\kappa)} I_0(r\kappa)$$

The log-likelihood is then:

$$\begin{aligned}
\ln(L(D_i; \kappa, \mu)) &= \log\left(\frac{1}{(2\pi)^{n-1}} \frac{1}{\prod_{i=1}^n I_0(\kappa)} I_0(r\kappa)\right) \\
\ln(L(\theta; \kappa, \mu)) &= \log\left(\frac{1}{(2\pi)^{n-1}}\right) + \log\left(\frac{1}{\prod_{i=1}^n I_0(\kappa)}\right) + \log(I_0(r\kappa)) \\
\ln(L(\theta; \kappa, \mu)) &= \log(1) - (n-1) \log((2\pi)) + \log(1) - \log\left(\sum_{i=1}^n I_0(\kappa)\right) + \log(I_0(r\kappa)) \\
\ln(L(\theta; \kappa)) &= -(n-1) \log((2\pi)) - \log\left(\sum_{i=1}^n I_0(\kappa)\right) + \log(I_0(r\kappa))
\end{aligned}$$

We construct the function `loglikelihood` in R to obtain the values from the last equation using the given data “Bees”:

```
loglikelihood <- function(x,kappa) {
  n <- length(x)
  c <- -(n-1)*log(2*pi)
  r <- sqrt((sum(cos(x)))^2 + (sum(sin(x)))^2)
  b_k <- besseli(kappa, 0)
  b_rk <- besseli(r*kappa, 0)
  loglike <- c - log(sum(b_k)) + log(b_rk)
  loglike
}
```

Based on the document in quercus we can compute the pre-normalized $U(\theta)$ as:

$$\mathfrak{U}(\theta) = \exp[\ln(\pi(\theta)) + \ln(L(\theta)) - \max_{\theta} \{ \ln(\pi(\theta)) + \ln(L(\theta)) \}]$$

We can define the prior here as:

$$\pi(\kappa, \mu) = \frac{\lambda}{2\pi} \exp(-\lambda \kappa)$$

So, our pre-normalized $U(\theta)$ has the following form:

$$\begin{aligned}
\mathfrak{U}(\theta) &= \exp \left[\ln\left(\frac{\lambda}{2\pi} \exp(-\lambda \kappa)\right) - (n-1) \log((2\pi)) - \log\left(\sum_{i=1}^n I_0(\kappa)\right) + \log(I_0(r\kappa)) \right. \\
&\quad \left. - \max_{\theta} \left\{ \ln\left(\frac{\lambda}{2\pi} \exp(-\lambda \kappa)\right) - (n-1) \log((2\pi)) - \log\left(\sum_{i=1}^n I_0(\kappa)\right) + \log(I_0(r\kappa)) \right\} \right]
\end{aligned}$$

Where the log prior has the following form:

$$\begin{aligned}
&\ln\left(\frac{\lambda}{2\pi} \exp(-\lambda \kappa)\right) \\
&\ln\left(\frac{\lambda}{2\pi}\right) + \ln(\exp(-\lambda \kappa)) \\
&\ln(\lambda) - \ln(2\pi) - \lambda \kappa
\end{aligned}$$

Then, we construct our function `prenorm` that calculates the pre-normalized using the log-prior, log-likelihood, and given parameters and data:

```

prenorm <- function(x, kappa, lambda) {
  a <- loglikelihood(x,kappa)
  ln_prior <- log(lambda) - log(2*pi) - lambda*kappa
  a <- a + ln_prior # add log-prior
  a <- a - max(a) # subtract maximum
  pre <- exp(a) # pre-normalized
  pre
}

```

However, this is not enough we still need to compute the value of kappa hat, note that a simple estimate of κ is:

$$\hat{\kappa} = \frac{\frac{r}{n}(2 - \frac{r^2}{n^2})}{1 - \frac{r^2}{n^2}}$$

Here r is defined as:

$$r = \left\{ \left(\sum_{i=1}^n \cos(d_i) \right)^2 + \left(\sum_{i=1}^n \sin(d_i) \right)^2 \right\}^{\frac{1}{2}}$$

Given all this information and the number of observations in bees (i.e. $n = 279$) we estimate κ :

$$\hat{\kappa} = \frac{\frac{\left\{ \left(\sum_{i=1}^n \cos(d_i) \right)^2 + \left(\sum_{i=1}^n \sin(d_i) \right)^2 \right\}^{\frac{1}{2}}}{279} \left(2 - \frac{\left\{ \left(\sum_{i=1}^n \cos(d_i) \right)^2 + \left(\sum_{i=1}^n \sin(d_i) \right)^2 \right\}}{279^2} \right)}{1 - \frac{\left\{ \left(\sum_{i=1}^n \cos(d_i) \right)^2 + \left(\sum_{i=1}^n \sin(d_i) \right)^2 \right\}}{279^2}}$$

Construct the function `kappa_hat` in order to compute the value that maximizes kappa:

```

kappa_hat <- function(x){
  n <- length(x)
  a <- sqrt((sum(cos(x)))^2 + (sum(sin(x)))^2)/n
  b <- sqrt((sum(cos(x)))^2 + (sum(sin(x)))^2)/(n^2)
  kappa_hat <- ( a *(2-b) ) / (1-b)
  kappa_hat
}

```

Using the given dataset:

```
kappa_hat(bees)
```

```
## [1] 0.0536232
```

So the posterior density should be the largest for values of κ close to 0.0536232. With this information we can set a range of values for κ in order to construct the posterior density.

In the following R chunk we summary the functions that we constructed:

```

loglikelihood <- function(x,kappa) {
  n <- length(x)
  c <- (n-1) * log(2*pi)
  r <- sqrt((sum(cos(x)))^2 + (sum(sin(x)))^2)
}

```

```

b_k <- besseli(kappa, 0)
b_rk <- besseli(r*kappa, 0)
loglike <- - c - log(sum(b_k)) + log(b_rk)
loglike
}

prenorm <- function(x, kappa, lambda) {
  a <- loglikelihood(x,kappa)
  ln_prior <- (log(lambda) - log(2*pi) - (lambda*kappa))
  a <- a + ln_prior # add log-prior
  a <- a - max(a) # subtract maximum
  pre <- exp(-a) # pre-normalized, minus to avoid problems with the numerical integration
  pre
}

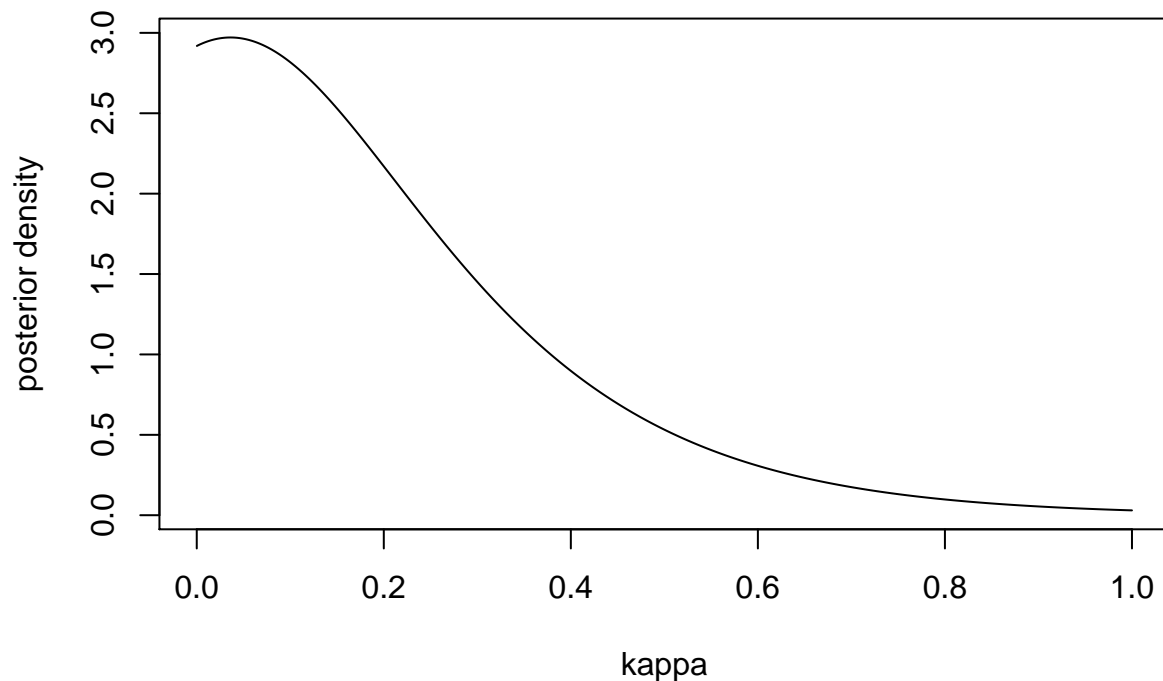
kappa_hat <- function(x){
  n <- length(x)
  a <- sqrt( (sum(cos(x)))^2 + (sum(sin(x)))^2 )/n
  b <- sqrt((sum(cos(x)))^2 + (sum(sin(x)))^2)/(n^2)
  kappa_hat <- ( a *(2-b) ) / (1-b)
  kappa_hat
}

```


```

kappa <- c(1:10000)/10000 # 10000 values of kappa
prenorm_post <- prenorm(bees, kappa = kappa , lambda = 1) # compute pre-normalized posterior
mult <- c(1/2,rep(1,9998),1/2) # multipliers for trapezoidal rule
norm <- sum(mult*prenorm_post)/10000 # integral evaluated using trapezoidal rule
post <- prenorm_post/norm # normalized posterior
plot(kappa,post,type="l",ylab="posterior density")

```



```
post.cdf <- cumsum(mult*post)/10000 # compute the posterior cdf
plot(kappa,post.cdf,type="l",ylab="cumulative posterior probability")
abline(h=c(0.025,0.975),lty=2)
```




```
lower <- max(kappa[post.cdf<0.025]) # lower limit for credible interval
upper <- min(kappa[post.cdf>0.975]) # upper limit for credible interval
c(lower,upper)
```


```
## [1] 0.0085 0.6989
```

Now we do the same but for $\Lambda = 0.1$:

```
kappa <- c(1:10000)/10000 # 671 values of kappa
prenorm_post <- prenorm(bees, kappa = kappa , lambda = 0.1) # compute pre-normalized posterior
mult <- c(1/2,rep(1,9998),1/2) # multipliers for trapezoidal rule
norm <- sum(mult*prenorm_post)/10000 # integral evaluated using trapezoidal rule
post <- prenorm_post/norm # normalized posterior
plot(kappa,post,type="l",ylab="posterior density")
```



```
post.cdf <- cumsum(mult*post)/10000 # compute the posterior cdf
plot(kappa,post.cdf,type="l",ylab="cumulative posterior probability")
abline(h=c(0.025,0.975),lty=2)
```



```
lower <- max(kappa[post.cdf<0.025]) # lower limit for credible interval
upper <- min(kappa[post.cdf>0.975]) # upper limit for credible interval
c(lower,upper)
```

```
## [1] 0.0071 0.6274
```

Discussion of the Results

We can see that the posterior density for $\lambda = 1$ approaches zero slower than the posterior density for $\lambda = 0.1$. As mentioned before the estimated value of *kappa* that maximizes the posterior (marginal) density for κ is close to 0.05 and we can see that the peak is constant at this point in the two posterior densities here calculated.

Moreover, the values of *kappa* and the construction of the posterior distribution here presented rule out the possibility of $\kappa = 0$ and the values of κ are strictly positive.

As a final remark, note that the range of values of κ here considered are $[0.0001, 1]$, which includes the expected peak at 0.05. The credible interval for the posterior cdf when $\lambda = 1$ is $[0.025, 0.975]$ and for the case when $\lambda = 0.1$ is $[0.0071, 0.6274]$ which again shows that our previous claim that the density for the case when $\lambda = 0.1$ approaches zero more rapidly than the case when $\lambda = 1$.

(d) The model that we've used to this point assumes that κ is strictly positive and thus rules out the case where $\kappa = 0$ (which implies that D_1, \dots, D_n are uniformly distributed). We can actually address the general model (i.e. $\kappa \geq 0$) by putting a prior probability θ on the model $\kappa = 0$ (which implies that we put a prior probability $1 - \theta$ on $\kappa > 0$). The prior distribution on the parameter space $\{(\kappa, \mu) : \kappa \geq 0, 0 \leq \mu < 2\pi\}$ has a point mass of θ at $\kappa = 0$ and the prior density over $\{(\kappa, \mu) : \kappa > 0, 0 \leq \mu < 2\pi\}$ is $(1 - \theta)\pi(\kappa, \mu)$. (This is a simple example of a “**spike-and-slab**” prior, which are used in Bayesian model selection; in this case, the “spike” refers to the prior point mass θ and the “slab” is the prior density on (κ, μ) for $\kappa > 0$.)

The posterior probability that $\kappa = 0$ is then

$$\text{Prob}(\kappa = 0 | d_1, \dots, d_n) = \frac{\theta(2\pi)^{-n}}{\theta(2\pi)^{-n} + (1 - \theta) \int_0^\infty \int_0^{2\pi} \mathcal{L}(\kappa, \mu) \pi(\kappa, \mu) d\mu d\kappa}$$

Using the prior $\pi(\kappa, \mu)$ in part (b) with $\lambda = 1$, evaluate $\text{Prob}(\kappa = 0 | d_1, \dots, d_n)$ for $\theta = 0.1, 0.2, 0.3, \dots, 0.9$. (This is easier than it looks as you've done much of the work in part (c).)

Address the general model (i.e $\kappa \geq 0$) by putting a prior probability on the model $\kappa = 0$ (this implies a prior probability on $1 - \theta$ on $\kappa > 0$). The prior density over $\{(\kappa, \mu) : \kappa > 0, 0 \leq \mu < 2\pi\}$ is $(1 - \theta)\pi(\kappa, \mu)$, where

$$\pi(\kappa, \mu) = \frac{\lambda}{2\pi} \exp(-\lambda \kappa)$$

this last prior is used with $\lambda = 1$ and becomes:

$$\pi(\kappa, \mu) = \frac{1}{2\pi} \exp(-1 \kappa)$$

So, the prior for this part is:

$$(1 - \theta) \pi(\kappa, \mu) \equiv (1 - \theta) \frac{1}{2\pi} \exp(-1 \kappa)$$

Note that the new log prior looks like:

$$\begin{aligned} & \ln((1 - \theta) \frac{1}{2\pi} \exp(-1 \kappa)) \\ & \ln(1 - \theta) + \ln(\frac{1}{2\pi}) + \ln(\exp(-1 \kappa)) \\ & \ln(1 - \theta) - \ln(2\pi) - 1 \kappa \end{aligned}$$

Using the code from part c:

```
loglikelihood2 <- function(x,kappa) {
  n <- length(x)
  c <- (n-1) * log(2*pi)
  r <- sqrt((sum(cos(x)))^2 + (sum(sin(x)))^2)
  b_k <- besseli(kappa, 0)
  b_rk <- besseli(r*kappa, 0)
```

```

loglike <- c - log(sum(b_k)) + log(b_rk)
loglike
}

prenorm2 <- function(x, kappa, lambda, theta) {
  a <- loglikelihood2(x,kappa)
  ln_prior <- log(1-theta) + (log(lambda) - log(2*pi) - (lambda*kappa))
  a <- a + ln_prior # add log-prior
  a <- a - max(a) # subtract maximum
  pre <- exp(-a) # pre-normalized, minus to avoid problems with the numerical integration
  pre
}

```


For all thetas

```

#Kappa = 0
theta <- seq(from = 0.1, to = 0.9, by = 0.1)
prenorm_post <- prenorm2(bees, kappa = 0, lambda = 1, theta = rep(theta,10000)) # compute pre-normalized
mult <- c(1/2,rep(1,9998),1/2) # multipliers for trapezoidal rule
norm <- sum(mult*prenorm_post)/10000 # integral evaluated using trapezoidal rule
post <- prenorm_post/norm # normalized posterior
hist(post)

```


Histogram of post



```

plot(rep(theta,10000),post,type="l",ylab="posterior density", xlab="theta = 0.1,..., 0.9")

```



(c) Consider the following model:

$$Y_i = \beta_0 + \beta_1 x_i + \varepsilon_i \quad (i = 1, \dots, n)$$

where $x_i = i/n$ and $\{\varepsilon_i\}$ are independent $\mathcal{N}(0,1)$ random variables. For least squares estimation, the variance of $\hat{\beta}_1$ tends to 0 like constant/ n as $n \rightarrow \infty$. For LMS estimation, $\text{Var}(\hat{\beta}_1) \approx \gamma/n^\alpha$ for some $\gamma > 0$ and $\alpha > 0$. In lecture, we claimed that $\alpha = 2/3$. The theoretical proof of this is very technical; however, it is possible to estimate α via simulation.

The idea here is very simple: We can compute $\hat{\beta}_1$ based on n observations and replicate this process M times, which results in M values of $\hat{\beta}_1$. We can then estimate $\text{Var}(\hat{\beta}_1)$ from these M values, in which case

$$\begin{aligned} \widehat{\text{Var}}(\hat{\beta}_1) &\approx \gamma/n^\alpha \\ \text{or } \ln(\widehat{\text{Var}}(\hat{\beta}_1)) &\approx \ln(\gamma) - \alpha \ln(n) \end{aligned}$$

Repeating this process for a range of sample sizes allows us estimate α .

Estimate α based on sample sizes $n = 50, 100, 500, 1000, 5000$.

Consider the following model

$$Y_i = \beta_0 + \beta_1 x_i + \varepsilon_i, \quad (i = 1, \dots, n)$$

where: $x_i = \frac{i}{n}$ and ε_i are independent $\mathcal{N}(0,1)$ random variables.

For LMS estimation

$$\text{Var}(\hat{\beta}_1) \approx \gamma/n^\alpha$$

for some $\gamma > 0$ and $\alpha > 0$. In class we claimed that $\alpha = 1/3$, here estimate α via simulation.

Step 1) compute $\hat{\beta}_1$ based on $n=50$ obs

Step 2) Replicate this process M times (obtaining M values of $\hat{\beta}_1$).

Step 3) Estimate

$$\text{Var}(\hat{\beta}_1)$$

from these M values in which case:

$$\begin{aligned} \widehat{\text{Var}}(\hat{\beta}_1) &\approx \frac{\gamma}{n^\alpha} \\ \text{or: } \ln(\widehat{\text{Var}}(\hat{\beta}_1)) &\approx \ln(\gamma) - \alpha \ln(n) \end{aligned}$$

Using a sample size of 50 (i.e: $n = 50$):

```
library(MASS)
n <- 50
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  y <- rnorm(n, 0, 1)
  r <- lmsreg(y~x)
```

```

    beta <- c(beta,r$coef[2])
  }
  beta_50 <- beta
  save(beta_50, file = "beta_50.RData")

```

```

load("beta_50.RData")
var50 <- var(beta_50)
var50

```

```
## [1] 1.388859
```

Using a sample size of 50 (i.e: $n = 100$):

```

library(MASS)
n <- 100
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  y <- rnorm(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta,r$coef[2])
}
beta_100 <- beta
save(beta_100, file = "beta_100.RData")

```

```

load("beta_100.RData")
var100 <- var(beta_100)
var100

```

```
## [1] 0.8161859
```

Using a sample size of 50 (i.e: $n = 500$):

```

library(MASS)
n <- 500
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  y <- rnorm(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta,r$coef[2])
}
beta_500 <- beta
save(beta_500, file = "beta_500.RData")

```

```

load("beta_500.RData")
var500 <- var(beta_500)
var500

```

```
## [1] 0.210944
```

Using a sample size of 50 (i.e: $n = 1000$):

```
library(MASS)
n <- 1000
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  options(mc.cores=parallel::detectCores())
  y <- rnorm(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta,r$coef[2])
}
beta_1000 <- beta
save(beta_1000, file = "beta_1000.RData")
```

```
load("beta_1000.RData")
var1000 <- var(beta_1000)
var1000
```

```
## [1] 0.1419492
```

Using a sample size of 50 (i.e: $n = 5000$):

```
library(MASS)
n <- 5000
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  options(mc.cores=parallel::detectCores())
  y <- rnorm(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta,r$coef[2])
}
beta_5000 <- beta
save(beta_5000, file = "beta_5000.RData")
```

```
load("beta_5000.RData")
var5000 <- var(beta_5000)
var5000
```

```
## [1] 0.04482628
```

Summarizing the variances obtained:

```
cbind(var50, var100, var500, var1000, var5000)
```

```
##          var50    var100    var500    var1000    var5000
## [1,] 1.388859 0.8161859 0.210944 0.1419492 0.04482628
```

Now we only need to estimate α , we can do this using a simple OLS estimation:

$$\ln(\hat{Var}(\beta_1)) \approx \ln(\gamma) - \alpha \ln(n)$$

So we take the logarithm of our estimated variances:

```
y <- log(c(var50, var100, var500, var1000, var5000))
log(rbind(c(var50, var100, var500, var1000, var5000)))
```

```
##           [,1]      [,2]      [,3]      [,4]      [,5]
## [1,] 0.3284828 -0.2031131 -1.556163 -1.952286 -3.104961
```

This is corresponding for each log-value of n, that is:

```
##           n50      n100      n500      n1000      n5000
## [1,] 3.912023 4.60517 6.214608 6.907755 8.517193
```

So, we run our regression as follows:

```
library(MASS)
ols_reg <- lm(y~x)
ols_reg <- ols_reg$coefficients
lms <- lmsreg(y~x)
lms_reg <- lms$coefficients
```

Note that we run the OLS regression and LMS regression just to comparison purposes. The estimates for the coefficients that these regressions produce are very similar:

```
##           ols_reg      lms_reg
## (Intercept) 3.221948 3.3049219
## x          -0.749344 -0.7613393
```

Finally, our estimate of α is **0.749** for the OLS regression and **0.7612** for the LMS regression, this is based on samples sizes of $n = 50, 100, 500, 1000, 5000$ and the assumption that the errors are independent $N(0,1)$ random variables.

(d) Repeat part (c) using Cauchy errors. (For Cauchy errors, the variance of the least squares estimator does not tend to 0 as $n \rightarrow \infty$.) Do you get a similar value of α ?

```
library(MASS)
n <- 50
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  y <- rcauchy(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta, r$coef[2])
}
beta_50_cauchy <- beta
save(beta_50_cauchy, file = "beta_50_cauchy.RData")
```

```
load("beta_50_cauchy.RData")
var50_cauchy <- var(beta_50_cauchy)
var50_cauchy
```

```
## [1] 1.292033
```

```
library(MASS)
n <- 100
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  y <- rcauchy(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta, r$coef[2])
}
beta_100_cauchy <- beta
save(beta_100_cauchy, file = "beta_100_cauchy.RData")
```

```
load("beta_100_cauchy.RData")
var100_cauchy <- var(beta_100_cauchy)
var100_cauchy
```

```
## [1] 0.663055
```

```
library(MASS)
n <- 500
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  y <- rcauchy(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta, r$coef[2])
}
```



```

}
beta_500_cauchy <- beta
save(beta_500_cauchy, file = "beta_500_cauchy.RData")

```

```

load("beta_500_cauchy.RData")
var500_cauchy <- var(beta_500_cauchy)
var500_cauchy

```

```
## [1] 0.1980633
```

```

library(MASS)
n <- 1000
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  options(mc.cores=parallel::detectCores())
  y <- rcauchy(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta,r$coef[2])
}
beta_1000_cauchy <- beta
save(beta_1000_cauchy, file = "beta_1000_cauchy.RData")

```

```

load("beta_1000_cauchy.RData")
var1000_cauchy <- var(beta_1000_cauchy)
var1000_cauchy

```

```
## [1] 0.1247424
```

```

library(MASS)
n <- 5000
nrep <- 1000
x <- c(1:n)/n
beta <- NULL
for (i in 1:nrep) {
  options(mc.cores=parallel::detectCores())
  y <- rcauchy(n, 0, 1)
  r <- lmsreg(y~x)
  beta <- c(beta,r$coef[2])
}
beta_5000_cauchy <- beta
save(beta_5000_cauchy, file = "beta_5000_cauchy.RData")

```

```

load("beta_5000_cauchy.RData")
var5000_cauchy <- var(beta_5000_cauchy)
var5000_cauchy

```

```
## [1] 0.04403548
```

We summarized the estimated variances for β_1 when using the cauchy errors:

```
cbind(var50_cauchy, var100_cauchy, var500_cauchy, var1000_cauchy, var5000_cauchy)

##      var50_cauchy var100_cauchy var500_cauchy var1000_cauchy var5000_cauchy
## [1,]      1.292033      0.663055      0.1980633      0.1247424      0.04403548
```

In log-terms we have:

```
y_c <- log(c(var50_cauchy, var100_cauchy, var500_cauchy, var1000_cauchy, var5000_cauchy))
log(cbind(var50_cauchy, var100_cauchy, var500_cauchy, var1000_cauchy, var5000_cauchy))

##      var50_cauchy var100_cauchy var500_cauchy var1000_cauchy var5000_cauchy
## [1,]      0.2562169      -0.4108974      -1.619169      -2.081505      -3.12276
```

And using the same approach as part c we have:

```
library(MASS)
ols_reg <- lm(y_c~x)
ols_reg_c <- ols_reg$coefficients
lms <- lmsreg(y_c~x)
lms_reg_c <- lms$coefficients
```

So, the estimated coefficients for α and γ are summarized in the following table:

```
##      ols_reg_c  lms_reg_c
## (Intercept)  2.9986972  2.4341235
## x           -0.7285798 -0.6530013
```

When the errors are assumed to follow a Cauchy distribution the estimates for α differ from the case when the errors are assumed to come from a independently $N(0,1)$ distribution. Then, in the case of the OLS regression the estimate for α is **0.7286** and when using the LMS regression the estimate for α is **0.653** which is clearly a slightly different estimate than when using the OLS regression. This is based on samples sizes of $n = 50, 100, 500, 1000, 5000$ and the assumption that the errors are independent $N(0,1)$ random variables.

So, we conclude that the values for α when using Cauchy errors are slightly different than part c, we remark that for Cauchy errors the variance of the OLS estimator does not tend to zero as n gets large, this explains the difference in the estimates.